

# NEURAL NETWORKS – METHOD SUMMARY

REGRESSION

CLASSIFICATION

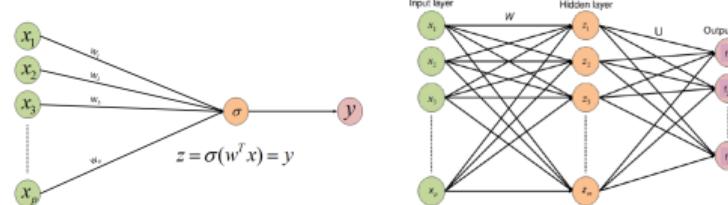
(NON)PARAMETRIC

BLACK-BOX

General idea

- Learn **composite function** through series of nonlinear feature transformations, represented as **neurons**, organized hierarchically in **layers**
  - Basic neuron operation: 1) affine **transformation**  $\phi$  (weighted sum of inputs), 2) nonlinear **activation**  $\sigma$
  - Combinations of simple building blocks to create a complex model
- Optimize via **mini-batch stochastic gradient descent (SGD)** variants:
  - Gradient of each weight can be inferred from the **computational graph** of the network  
→ **Automatic Differentiation** (AutoDiff)
  - Algorithm to compute weight updates based on the loss is called **Backpropagation**

Hypothesis space  $\mathcal{H} = \left\{ f(\mathbf{x}) : f(\mathbf{x}) = \tau \circ \phi \circ \sigma^{(h)} \circ \phi^{(h)} \circ \sigma^{(h-1)} \circ \phi^{(h-1)} \circ \dots \circ \sigma^{(1)} \circ \phi^{(1)}(\mathbf{x}) \right\}$



# NEURAL NETWORKS – METHOD SUMMARY

## Architecture

- Input layer: original features  $\mathbf{x}$
- Hidden layers: nonlinear transformation of previous layer  $\phi^{(h)} = \sigma^{(h-1)}(\phi^{(h-1)})$
- Output layer: number of output neurons and activation depends on problem  $\tau(\phi)$ 
  - Regression: one output neuron,  $\tau = \text{identity}$
  - Binary classification: one output neuron,  $\tau = \frac{1}{1+\exp(-\theta^\top \mathbf{x})}$  (logistic sigmoid)
  - Multiclass Classification:  $g$  output neurons,  $\tau_j = \frac{\exp(f_j)}{\sum_{j=1}^g \exp(f_j)}$  (softmax)

Empirical risk In general, compatible with any differentiable loss

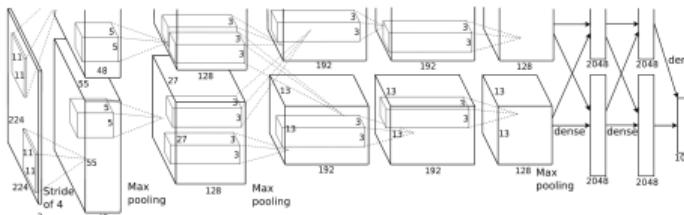
## Optimization

- Variety of different optimizers, mostly based on some form of **stochastic gradient descent (SGD)**
- Improvements:
  - (1) Accumulation of previous gradients → **Momentum**
  - (2) Weight specific scaling based on previous squared gradients → **RMSProb**  
⇒ **ADAM** combines (1) and (2)
  - (3) Learning rate schedules, e.g., decaying or cyclical learning rates
- Training progress is measured in full passes over the full training data, called **epochs**
- **Batch size** is a hyperparameter and limited by input data dimension

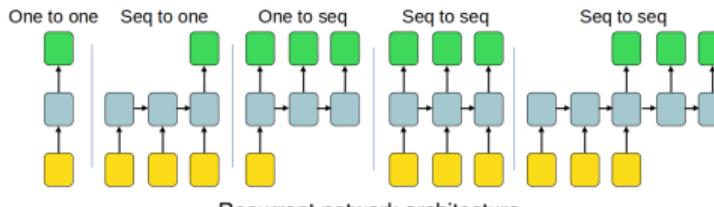
# NEURAL NETWORKS – METHOD SUMMARY

Network types Large variety of architectures for different data modalities

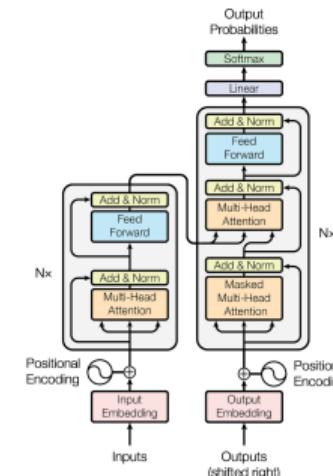
- **Feedforward NNs / multi-layer perceptrons (MLPs)**: sequence of **fully-connected** layers ⇒ tabular data
- **Convolutional NNs (CNNs)**: sequence of feature map extractors with spatial awareness ⇒ images, time series
- **Recurrent NNs (RNNs)**: handling of sequential, variable-length information ⇒ times series, text, audio
- **Transformers**: Learning invariances from data, handling multiple/any data modalities



Convolutional network architecture



Recurrent network architecture



Transformer network architecture

# NEURAL NETWORKS – METHOD SUMMARY

## Hyperparameters

- **Architecture:**

- Lots of design choices ⇒ tuning problem of its own.
- Typically: hierarchical optimization of components (cells) and macro structure of network  
→ **Neural Architecture Search (NAS)**
- Many predefined (well working) architectures exist for standard tasks

- **Training:**

- Initial learning rate and various regularization parameters
- Number of epochs is determined by **early-stopping**
- **Data-augmentation**, e.g., applying random rotations to input images

## Foundation models

- **Enormous** models trained on vast amounts of (general) data, e.g., all of wikipedia, in **self-supervised** fashion
- Used as starting point (**pre-trained**) and fine-tuned via **transfer** or **few-shot** learning for other tasks

## General hints

- Instead of NAS, use a standard architecture and tune training hyperparameters
- Training pipeline (data-augmentation, training schedules, ...) is more crucial than the specific architecture
- While NNets are state-of-the-art for **computer vision (CV)** and **natural language processing (NLP)**, we recommend not to use them for tabular data because alternatives perform better

## Disadvantages

- Typically, high computational **cost**
- High demand for **training data**
- Strong tendency to **overfit**
- Requiring lots of **tuning expertise**
- **Black-box** model – hard to interpret

## Advantages

- + Applicable to **complex, nonlinear** problems
- + Very **versatile** w.r.t. architectures
- + Strong **performance** if done right
- + Built-in **feature extraction**, obtained by intermediate representations
- + Easy handling of **high-dimensional** data
- + **Parallelizable** training